

AI Act Complete Pre-Consultory Report

Meridian Retail Group

Illustrative mid-size EU e-commerce retailer, 180 staff, sells in 6 EU markets

REFERENCE	STATUS	UPDATED
SPECIMEN	ILLUSTRATIVE	2026-07-26

Published so you can inspect the format, depth and decision logic before buying. Meridian Retail Group is invented; no real client information appears in this document.

How to read this specimen

This is the full structure of the paid RapidAct assessment, filled with an invented company's systems. The report separates observed facts, classification, required actions and issues that need specialist legal advice. It does not present the website, scanner or badge as evidence of compliance.

Document	RapidAct specimen assessment
Prepared for	Meridian Retail Group
Prepared by	Agents AI Ltd. · RapidAct
Status	Illustrative / not legal advice
Reference	SPECIMEN

CONTENTS

01. Your AI inventory	3
02. Risk classification, and why each system lands where it does	4
03. What you must disclose, where, and in what words	5
04. Your evidence position	6
05. Prioritised actions	7
06. Assessment	8

01. Your AI inventory

Every AI system the company operates, written down in one place. Most companies have never seen this list, and it is common for the item carrying the most exposure to be one nobody thought of as AI.

System	Where it runs	Vendor	Classification
Customer support chatbot	All storefront pages, bottom-right widget	Intercom Fin	Transparency, Art. 50(1)
AI-written product descriptions	~4,100 product detail pages	In-house, GPT-based pipeline	Transparency, Art. 50(2)
Recommendation and ranking engine	Home, category and basket pages	In-house	Out of scope for Art. 50
CV screening and ranking	Careers portal, applicant shortlisting	Third-party ATS module	High-risk, Annex III(4)(a)
AI voice agent for order status	Customer service phone line	Third-party voice platform	Transparency, Art. 50(1)
Internal code assistant	Engineering team only	GitHub Copilot	Out of scope

02. Risk classification, and why each system lands where it does

The support chatbot and the voice agent both interact directly with natural persons, which brings them within Article 50(1). The obligation is to ensure the person is informed that they are interacting with an AI system, unless that is obvious to a reasonably well-informed user. A branded widget labelled only with a first name does not meet that bar, and a voice agent that opens with a human-sounding greeting clearly does not.

The AI-written product descriptions engage Article 50(2), which concerns synthetic content and its marking in a machine-readable format. This is the obligation companies most often miss, because the output does not feel like "AI content" once it has been edited by a person. Where a human genuinely reviews and takes editorial responsibility for the text, the analysis changes; this section sets out where that line sits for your workflow.

The recommendation engine is assessed as out of scope for Article 50. It does not interact conversationally and does not generate synthetic media. It remains relevant to other regimes, notably the GDPR provisions on automated decision-making and the platform rules on ranking transparency, which are flagged but not assessed here.

CV screening is the system carrying the most exposure. Recruitment and applicant selection fall under Annex III(4)(a), which makes it high-risk rather than transparency-only. That is a materially heavier set of obligations covering risk management, data governance, logging, human oversight and registration. It also runs on a different clock.

The internal code assistant is out of scope. It does not touch external users and does not produce content published to the public. It is listed for completeness, because an inventory that omits systems is not an inventory.

The two clocks. Article 50 transparency duties apply from 2 August 2026. The Annex III high-risk obligations that catch your CV screening were postponed to December 2027. Both apply to this company; they simply do not start together. Treating the postponement as covering everything is the most expensive misreading we encounter.

03. What you must disclose, where, and in what words

For the support chatbot, disclosure must be available before the visitor sends their first message, not in a policy page and not after the conversation has begun. In practice this means a persistent line in the widget header or an opening message that states it plainly. Suggested wording: "You are chatting with an AI assistant. Ask for a human at any time and we will transfer you."

For the voice agent, the disclosure must come at the start of the call, in the same language as the call, before any question is asked of the caller. Suggested wording: "This call is handled by an automated AI assistant. Say 'agent' at any point to reach a person."

For AI-written product descriptions, the marking obligation is machine-readable, so a visible badge alone is insufficient. This is satisfied at generation time by embedding provenance metadata in the content pipeline rather than by editing 4,100 pages by hand. The implementation options are set out below in order of effort.

Because Meridian sells into six EU markets, each disclosure must be present in the language in which the service is offered. A single English notice on a Spanish-language storefront does not discharge the duty.

04. Your evidence position

The duty is not only to disclose but to be able to demonstrate that you disclosed. A regulator, an enterprise customer running vendor due diligence, or an insurer assessing a claim can each ask you to show what was displayed, on which pages, and from what date.

As described, Meridian can currently produce screenshots taken on the day of the request. That evidences the present, not the period under question, and is the single most common gap we find. What is needed is a dated record of disclosure events that predates the question being asked.

For the high-risk CV screening system the standard is considerably higher, extending to logs of system operation, documentation of the human oversight actually exercised, and records of the data used. If the ATS vendor holds these, you need a contractual route to obtain them, because the obligation sits with you as the deployer and cannot be delegated by silence.

Practical consequence. Evidence is cheap to start and impossible to backfill. A log begun today covers the period from today. A log begun when the question arrives covers nothing that matters.

05. Prioritised actions

Ordered by deadline and severity, separating what is genuinely urgent from what is merely being sold as urgent.

#	Action	Owner	Deadline
1	Open the documentation and logging question with your ATS vendor in writing	Procurement / HR	Start now, long lead time
2	Add AI disclosure to the support chatbot, in all six market languages	Web team	Before 2 August 2026
3	Add an opening AI disclosure to the voice agent script	Customer service	Before 2 August 2026
4	Begin a dated disclosure evidence log across the storefront	Web team	Immediately, it cannot be backfilled
5	Document your editorial review position for AI-written product copy	E-commerce / Legal	Before 2 August 2026
6	Add machine-readable provenance marking to the description pipeline	Engineering	Before 2 August 2026
7	Prepare the Annex III high-risk compliance programme for CV screening	HR / Legal	Before December 2027

06. Assessment

Meridian's exposure is concentrated in two places, and they are not the two places the team expected. The chatbot, which prompted the enquiry, is a straightforward fix: a wording change and a disclosure component, achievable in an afternoon at no licence cost.

The real exposure is the CV screening system, which nobody flagged because it was procured as a recruitment feature rather than as AI. It is high-risk under Annex III, it sits with a third-party vendor, and the contractual position on logs and documentation is currently unresolved. That should be opened with the vendor now, because it is the item with the longest lead time and the one you cannot fix unilaterally in the final week.

The product-description pipeline sits in between. It is genuinely arguable in either direction depending on how much editorial control your team exercises, and it is the one item here where we would recommend a documented internal position and, if the volume grows, a short legal opinion. What to put in front of counsel is set out above so that conversation is a narrow question rather than an open brief.

Nothing in this assessment requires you to stop using any system you currently operate.

Scope note. This document is a technical and organisational compliance assessment. It is not legal advice. Where a question turns on a point of law, the report identifies that question for counsel.